

Dynamic Programming and Applications

Discrete Time Dynamics and Optimization

Lecture 2

Andreas Schaab

Last lecture, we introduced the Bellman equation:

$$V(k) = \max_{k'} \left\{ u\left(f(k) - k'\right) + \beta V(k') \right\}$$

- The value today is = flow payoff + continuation value
- This assumes there is no uncertainty: We can make perfect forecasts about the future
Not how the world works!
- Suppose production also depends on productivity $y_t = f(k_t, A_t)$ and A_{t+1} is uncertain

In an uncertain world, we care about the *expected* continuation value

$$V(k, A) = \max_{k'} \left\{ u\left(f(k, A) - k'\right) + \beta \mathbb{E} V(k', A') \right\}$$

- Now the only question is: How do we compute the expectation \mathbb{E} ?
- We have to study stochastic processes (and stochastic calculus) to answer this

Outline

Part 1: Difference equations

1. Stochastic processes
2. Markov chains
3. Difference equations
4. Stochastic difference equations

Part 2: Stochastic dynamic programming

1. Stochastic dynamic programming
2. History notation
3. The stochastic neoclassical growth model

Part 3: Optimal stopping

Part 1: Difference Equations

1. Stochastic processes

- Let X_t be a random variable that is time t adapted
- Discrete time: We index time discretely $t = 0, 1, 2, \dots, T \leq \infty$
- Stochastic process in discrete time: a sequence of random variables indexed by t , $\{X_t\}_{t=0}^T$
- Continuous time: We index time continuously $t \in [0, T]$ with $T \leq \infty$
- Stochastic process in continuous time: a sequence of random variables indexed by t , $\{X_t\}_{t \geq 0}$

2. Markov chains

- A stochastic process $\{X_t\}$ has the *Markov property* if for all $k \geq 1$ and all t :

$$\mathbb{P}(X_{t+1} = x \mid X_t, X_{t-1}, \dots, X_{t-k}) = \mathbb{P}(X_{t+1} = x \mid X_t)$$

- *State space* of the Markov process = set of events or states that it visits
- A Markov chain is a Markov process (stochastic process with Markov property) that visits a finite number of states (*discrete state space*)
- Simplest example: Individual i is randomly hit by earnings (employment) shocks and switches between $X_t \in \{X^L, X^H\}$

- Markov chains have a *transition matrix* P that describes the probability of transitioning from state i to state j
- Simplest example with state space $\{X^L, X^H\}$

$$P = \begin{pmatrix} 0.8 & 0.2 \\ 0.2 & 0.8 \end{pmatrix}$$

- This says: P of staying in employment state = 0.8, P of switching = 0.2
- P_{ij} is the probability of switching from state i to state j (one period)
- P^2 characterizes transitions over two periods: $(P^2)_{ij}$ is prob of going from i to j in two periods
- The rows of the transition matrix have to sum to 1 (definition of probability measure)

3. Difference equations

- We start with deterministic (non-random) dynamics and then conclude with stochastic (random) dynamics
- The *first-order linear difference equation* is defined by

$$x_{t+1} = bx_t + cz_t \tag{1}$$

where $\{z_t\}$ is an exogenously given, bounded sequence

- For now, all objects are (real) scalars (easy to extend to vectors and matrices)
- Suppose we have an *initial condition* (i.e., given initial value) x_0
- When $c = 0$, (1) is a *time-homogeneous* difference equation
- When cz_t is constant for all t , (1) is an *autonomous* difference equation

Autonomous equations

- Consider the autonomous equation with $z_t = 1$
- A particular solution is the constant solution with $x_t = \frac{c}{1-b}$ when $b \neq 1$
- Such a point is called a *stationary point* or *steady state*
- General solution of the autonomous equation (for some constant x):

$$x_t = (x_0 - x)b^t + x \quad (2)$$

- Important question is long-run behavior (stability / convergence)
- When $|b| < 1$, (2) converges asymptotically to steady state x for any initial value x_0 (steady state x is globally stable)
- If $|b| > 1$, (2) explodes and is not stable (except when $x_0 = x$)

Examples in macro

Capital accumulation:

$$K_{t+1} = (1 - \delta)K_t + I_t$$

- δ is depreciation and I_t is investment
- This is a *forward equation* and requires an initial condition K_0
- If $I_t = 0$ and $0 < \delta < 1$, $K_t \rightarrow 0$
- If $I_t = c$ constant, then K_t converges to $\frac{c}{\delta}$: $K_{t+1} = (1 - \delta)\frac{c}{\delta} + c = \frac{c}{\delta}$

Wealth dynamics:

$$a_{t+1} = R_t a_t + y_t - c_t$$

- R_t is the gross real interest rate, y_t is income, c_t is consumption
- This is a *forward equation* and requires an initial condition a_0
- We will study this as a *controlled* process because c_t will be chosen optimally
- Work out the following: $R_t = R$ and $y_t = y$ constant, and

$$c_t = \left(1 - \frac{1}{R}\right) \left(a_t + \sum_{s=t}^{\infty} R^{-(s-t)} y\right)$$

What are the dynamics of a_t ?

Consumption Euler equation:

$$\frac{1}{C_t} = \beta R_t \frac{1}{C_{t+1}}$$

- $\frac{1}{C_t} = u'(C_t)$ is marginal utility with log preferences
- This is a *backward equation* and requires a terminal condition or transversality condition, i.e., c_T must converge to something
- Suppose there exists time T s.t. for all $t \geq T$, $C_t = C$
- Then solve *backwards* from: $\frac{1}{C_{T-1}} = \beta R_{T-1} \frac{1}{C_T}$ or expressed as *time-homogeneous first-order linear difference equation*

$$C_{T-1} = \frac{1}{\beta R_{T-1}} C_T$$

- Difference between *forward* and *backward* equations is critical! This is closely related to the idea of *boundary conditions* (much more to come)

4. Stochastic difference equations

- Consider the process $\{X_t\}$ with

$$X_{t+1} = AX_t + Cw_{t+1} \quad (3)$$

where w_{t+1} is an iid. process with $w_{t+1} \sim \mathcal{N}(0, 1)$

- Equation (3) is a *first-order, linear stochastic difference equation*
- Let \mathbb{E}_t the *conditional expectation* operator (conditional on time t information)
- For example:

$$\begin{aligned} \mathbb{E}_t(X_{t+1}) &= \mathbb{E}(X_{t+1} \mid X_t) = \mathbb{E}(AX_t + Cw_{t+1} \mid X_t) \\ &= AX_t + C\mathbb{E}(w_{t+1} \mid X_t) = AX_t + C\mathbb{E}(w_{t+1}) = AX_t \end{aligned}$$

- Rational expectations: agents' beliefs about stochastic processes are consistent with the true distribution of the process
- Key equation: wealth dynamics with income fluctuations:

$$a_{t+1} = R_t a_t + y_t - c_t,$$

where y_t is a stochastic process

- Consumption Euler equation with uncertainty (e.g., stochastic income):

$$u'(C_t) = \beta R \mathbb{E}_t \left[u'(C_{t+1}) \right]$$

Part 2: Stochastic Dynamic Programming

1. Stochastic dynamic programming

- Follow Ljungqvist-Sargent notation, Chapter 3.2
- Under uncertainty, household problem takes the form

$$\max_{\{c_t\}} \mathbb{E}_0 \sum_{t=0}^{\infty} \beta^t u(c_t)$$

subject to $k_{t+1} = g(k_t, c_t, \epsilon_{t+1})$ (*first-order stochastic difference equation*)

- Notice: $\{c_t\}$ now denotes a stochastic process, no longer simple sequence!
- $\{\epsilon_t\}_{t=0}^{\infty}$ is sequence of iid random variables (*stochastic process*)
- Initial condition x_0 given

- Usually best to start with sequence problem, then derive recursive representation
- To derive recursive representation, your first question must be:
 - Recursive representation means we go from thinking about sequences (stochastic processes) to thinking about functions
 - But functions *of what?* I.e., what is the domain of the functions we're interested in?
 - Answer: functions of the *state variables*
- What are state variables?
 - In the (deterministic) neoclassical growth model: just k
 - Generally: state variables = set of information you need today to compute the continuation value for tomorrow
 - That's why they're called "states"

- Dynamic programming: look for recursive representation with state variable k
- Q: Why is k the state variable here, not (k, ϵ) ? (Think about structure of $g(\cdot)$.)
- The problem is to look for a *policy function* $c(k)$ that solves

$$V(k) = \max_c \left\{ u(c) + \beta \mathbb{E} \left[V(g(k, c, \epsilon)) \mid k \right] \right\}, \quad \text{where } \mathbb{E}[V(\cdot) \mid k] = \int V(\cdot) dF(\epsilon)$$

- $V(k)$ is (lifetime) value that agent obtains from solving this problem starting from k
- FOC that characterizes the consumption policy function $c(k)$ is

$$0 = u'(c(k)) + \beta \mathbb{E} \left\{ \partial_k V(g(k, c(k), \epsilon)) \cdot \partial_c g(k, c(k), \epsilon) \mid k \right\} = 0$$

2. History notation

- A very popular approach to deal with uncertainty in macro is to use history notation (Ljungqvist-Sargent, e.g., chapters 8, 12)
- Time is discrete and indexed by $t = 0, 1, \dots$
- At every t , there is a realization of a stochastic event $s_t \in \mathcal{S}$
- We denote the **history** of such events up to t by $s^t = \{s_0, s_1, \dots, s_t\}$
- The unconditional probability of history s^t is given by $\pi_t(s^t \mid s_0)$
- If Markov, $\pi_t(s^t \mid s_0) = \pi(s_t \mid s_{t-1})\pi(s_{t-1} \mid s_{t-2}) \dots \pi(s_0)$

- Crucial to understand notation:
 - $\{c_t\}_{t \geq 0}$ is the stochastic process
 - c_t is the random variable
 - $c_t(s^t)$ is the realization of the random variable at date t in history s^t
- The **lifetime value** of representative household is then defined as

$$V(s_0) = \sum_{t=0}^T \beta^t \sum_{s^t} \pi_t(s^t | s_0) u(c_t(s^t), \ell_t(s^t))$$

- Here we also allow household to choose labor supply ℓ_t
- *Generalizations*: heterogeneous beliefs, general preferences (Epstein-Zin), recursive formulation, multiple commodities, intergenerational considerations

3. Stochastic Growth Model

- Discrete time: $t \in \{0, 1, \dots, T\}$, where $T \leq \infty$
- At t , event $s_t \in \mathcal{S}$ is realized; history $s^t = (s_0, \dots, s_t)$ has probability $\pi_t(s^t)$
- Representative household has preferences over consumption $c_t(s^t)$ and labor $\ell_t(s^t)$

$$\sum_{t=0}^{\infty} \beta^t \sum_{s^t} \pi_t(s^t) u(c_t(s^t), \ell_t(s^t))$$

- Inada conditions $\lim_{c \rightarrow 0} u_c(c, \ell) = \lim_{\ell \rightarrow 0} u_\ell(c, \ell) = \infty$
- At $t = 0$, household endowed with k_0

- Technology, capital accumulation, and budget / resource constraint:

$$c_t(s^t) + i_t(s^t) = A_t(s^t)F(k_t(s^{t-1}), \ell_t(s^t))$$

$$k_{t+1}(s^t) = (1 - \delta)k_t(s^{t-1}) + i_t(s^t)$$

- $F(\cdot)$ is twice continuously differentiable and constant returns to scale
- Source of uncertainty is stochastic process for TFP $A_t(s^t)$
- Standard regularity conditions on $F(\cdot)$ (see LS)

Lagrangian approach to sequence problem

$$L = \sum_{t=0}^{\infty} \sum_{s^t} \beta^t \pi_t(s^t) \left\{ u(c_t(s^t), l_t(s^t)) \right. \\ \left. + \lambda_t(s^t) \left[A_t(s^t) F(k_t(s^{t-1}), \ell_t(s^t)) - c_t(s^t) + (1 - \delta)k_t(s^{t-1}) - k_{t+1}(s^t) \right] \right\}$$

- FOCs for $c_t(s^t)$, $\ell_t(s^t)$ and $k_{t+1}(s^t)$ are given by

$$u_c(s^t) = \lambda_t(s^t)$$

$$u_\ell(s^t) = u_c(s^t) A_t(s^t) F_\ell(s^t)$$

$$u_c(s^t) \pi_t(s^t) = \beta \sum_{s^{t+1}|s^t} u_c(s^{t+1}) \pi_{t+1}(s^{t+1}) \left[A_{t+1}(s^{t+1}) F_k(s^{t+1}) + (1 - \delta) \right]$$

- Summation over $(s^{t+1} | s^t)$ is like conditional expectation
summing over histories that branch out from s^t

Dynamic programming approach

- Assume time-homogeneous Markov process:

$$\mathbb{E}_t(A_{t+1}) = \mathbb{E}\left[A(s^{t+1}) \mid s^t\right] = \mathbb{E}\left[A(s_{t+1}) \mid s_t\right] = \sum_{s'} \pi(s' \mid s_t) A(s')$$

- Drop t subscripts: s is current state, s' denotes next period's draw
- Denote by X_t *endogenous state* (assume for now there is such a representation)
- Intuitively: s is the exogenous state and X is the endogenous state

Bellman equation becomes:

$$V(X, s) = \max_{c, \ell} \left\{ u(c, \ell) + \beta \sum_{s'} \pi(s' \mid s) V(X', s') \right\} \quad \text{where } X' = g(X, c, \ell, s, s')$$

Part 3: Optimal Stopping

Application: optimal stopping problem

Problem: Every period t , an agent draws an offer x from a uniform distribution over the unit interval $[0, 1]$. The agent can accept the offer, in which case her payoff is x , and the game ends, or the agent can reject the offer and draw again a period later. Draws are independent. Rejections are costly because the agent discounts the future at β . The game continues until the agent receives an offer she accepts.

Many applications (problems in life) look like this:

- buying a house
- searching for a partner
- closing a production plant
- exercising an option
- adopting a new technology

What is recursive / dynamic programming representation of optimal stopping problem?

What is recursive / dynamic programming representation of optimal stopping problem?

Agent's dynamic optimization problem given recursively by Bellman equation

$$V(x) = \max \left\{ x, \beta \mathbb{E} V(x') \right\}$$

where the expectation (operator) \mathbb{E} is taken over the next draw x'

Our problem is to find the value function $V(x)$ that solves the Bellman equation. We'll also want to find the associated policy rule.

Definition: *A policy is a function that maps every point in state space $[0, 1]$ to an action*

There are 2 actions: ACCEPT and REJECT

Definition: *An optimal policy achieves payoff $V(x)$ for all feasible $x \in [0, 1]$*

Let's try to understand the shape of $V(x)$ intuitively:

- For large values \hat{x} where you ACCEPT, what's the value $V(\hat{x})$?
- For small values \tilde{x} where you REJECT and instead choose the continuation value, $\beta \mathbb{E}V(x') > \tilde{x}$, does the continuation value depend on \tilde{x} ? Why not?

Let's try to understand the shape of $V(x)$ intuitively:

- For large values \hat{x} where you ACCEPT, what's the value $V(\hat{x})$?
- For small values \tilde{x} where you REJECT and instead choose the continuation value, $\beta \mathbb{E}V(x') > \tilde{x}$, does the continuation value depend on \tilde{x} ? Why not?

Shape of $V(x)$ must therefore be:

$$V(x) = \begin{cases} x & \text{if } x \geq x^* \\ x^* & \text{if } x < x^* \end{cases}$$

- Solution to the problem: there is a threshold $x^* \in [0, 1]$ s.t. agent accepts for $x \geq x^*$
- Also called **free boundary problem** (have to find endogenous boundary x^*)

Lemma: *In the optimal stopping problem, a policy is a best response to a continuation value function $\widehat{v}(x)$ if and only if the policy is a threshold rule with cutoff*

$$x^* \equiv \beta \mathbb{E}[\widehat{v}(x')]$$

Proof: Show by contradiction that optimization must imply

$$\begin{array}{ll} \text{ACCEPT} & \text{if } x > \beta \mathbb{E}[\widehat{v}(x')] \equiv x^* \\ \text{REJECT} & \text{if } x < \beta \mathbb{E}[\widehat{v}(x')] \equiv x^* \end{array}$$

If $x = \beta \mathbb{E}[\widehat{v}(x')]$, then ACCEPT and REJECT generate the same payoff. ■

- Why no jump in $V(x)$ at x^* ? (lim from RHS must be x^* , from LHS by contradiction)
- Continuation value must be $V(x')$ because problem tmr is repeat of today

- We just concluded: at $x = x^*$, indifferent between ACCEPT and REJECT
- This is enough information to solve the problem!

$$\begin{aligned} V(x^*) &= x^* \\ &= \beta \mathbb{E} V(x') \\ &= \beta \int_0^{x^*} x^* f(x) dx + \beta \int_{x^*}^1 x f(x) dx \\ &= \beta x^* [x]_0^{x^*} + \beta \frac{1}{2} [x^2]_{x^*}^1 \\ &= \beta (x^*)^2 + \beta \frac{1}{2} [1 - (x^*)^2] \end{aligned}$$

Solution:

$$x^* = \frac{1}{\beta} (1 - \sqrt{1 - \beta^2})$$

Always sanity-check comparative statics: What happens as $\beta \rightarrow 0$ and $\beta \rightarrow 1$?

Why is this threshold rule a *solution to the Bellman Equation*?

If you REJECT, your continuation payoff is

$$x^* = \beta \mathbb{E} V(x') = \beta \int_0^{x^*} x^* f(x) dx + \beta \int_{x^*}^1 x f(x) dx.$$

So it's optimal to REJECT if $x \leq x^*$ and it's optimal to ACCEPT if $x \geq x^*$. Hence, for all values of x

$$V(x) = \max\{x, x^*\} = \max\{x, \beta \mathbb{E}[V(x')]\}$$

Application: job market paper

Problem: Imagine you were a graduate student in economics. And you have to find your job market paper. Time is discrete, and you have to go on the job market in period T .

Application: job market paper

Problem: Imagine you were a graduate student in economics. And you have to find your job market paper. Time is discrete, and you have to go on the job market in period T .

Your current project has quality x . Each period, you allocate your time to (i) work on your current project or (ii) search for a new project. Your project quality evolves as:

Application: job market paper

Problem: Imagine you were a graduate student in economics. And you have to find your job market paper. Time is discrete, and you have to go on the job market in period T .

Your current project has quality x . Each period, you allocate your time to (i) work on your current project or (ii) search for a new project. Your project quality evolves as:

$$x' = \begin{cases} f(x, \epsilon) & \text{where } \epsilon \sim \log \mathcal{N}(\mu, \sigma) \\ \epsilon^{new} & \text{where } \epsilon^{new} \sim \log \mathcal{N}(\mu^{new}, \sigma^{new}) \end{cases}$$

Your goal is to go on the market with the highest-quality project.

Application: job market paper

Problem: Imagine you were a graduate student in economics. And you have to find your job market paper. Time is discrete, and you have to go on the job market in period T .

Your current project has quality x . Each period, you allocate your time to (i) work on your current project or (ii) search for a new project. Your project quality evolves as:

$$x' = \begin{cases} f(x, \epsilon) & \text{where } \epsilon \sim \log \mathcal{N}(\mu, \sigma) \\ \epsilon^{new} & \text{where } \epsilon^{new} \sim \log \mathcal{N}(\mu^{new}, \sigma^{new}) \end{cases}$$

Your goal is to go on the market with the highest-quality project.

What is the Bellman equation for this problem?

- This is an example of **non-stationary dynamic programming**
- Let's start at the end: What is your payoff when you go on the market in period T ?

$$V_T(x) = x$$

- What is your value in period $T - 1$?

$$\begin{aligned} V_{T-1}(x) &= \max \left\{ \mathbb{E} V_T \left(f(x, \epsilon) \right), \mathbb{E} V_T(\epsilon^{new}) \right\} \\ &= \max \left\{ \mathbb{E} f(x, \epsilon), \mu^{new} \right\} \end{aligned}$$

- Iterate backwards until

$$V_0(x) = \max \left\{ \mathbb{E} V_1 \left(f(x, \epsilon) \right), \mathbb{E} V_1(\epsilon^{new}) \right\}$$

- This is an easy problem to solve on the computer! (More examples to come in class)
You can probably solve this by hand but I haven't tried. Give it a shot for simple $f(\cdot)$!
- Interpret the economics:
 - How to think about the length of the time step? (How large is T ?)
 - How to think about $f(x, \epsilon)$?
 - Is σ^{new} larger or smaller than σ ?
 - Should σ maybe also be a state variable (and change over time)?
- Key takeaway from this sort of problem: $\sigma^{new} \gg 0$ is your best friend and much more important than μ^{new} !!! (Easy to check on computer!)
- While T is still large, look for the riskiest projects you can think of!!!